

大模型：分布式训练

PTD 并行配置



ZOMI

大模型业务全流程



大模型系列 – 分布式训练加速

- 具体内容

- I. 分布式加速库 :

- 业界常用分布式加速库 & 作用

2. DeepSpeed 特性 :

- 基本概念 - 整体框架 – Zero-1/2/3 – ZeRO-Offload – ZeRO-Infinity

3. Megatron 特性 :

- I. 总体介绍 – 整体流程 – 并行配置 – DP – TP – PP

3.2 Megatron-LM

PTD回顾与初始化

initialize_model_parallel() 并行配置初始化

```
12 # Intra-layer model parallel group that the current rank belongs to.  
13 _TENSOR_MODEL_PARALLEL_GROUP = None  
14 # Inter-layer model parallel group that the current rank belongs to.  
15 _PIPELINE_MODEL_PARALLEL_GROUP = None  
16 # Model parallel group (both intra- and pipeline) that the current rank  
17 _MODEL_PARALLEL_GROUP = None  
18 # Embedding group.  
19 _EMBEDDING_GROUP = None  
20 # Position embedding group.  
21 _POSITION_EMBEDDING_GROUP = None  
22 # Data parallel group that the current rank belongs to.  
23 _DATA_PARALLEL_GROUP = None  
24 _DATA_PARALLEL_GROUP_GLOO = None  
25 # tensor model parallel group and data parallel group combined  
26 # used for fp8 and moe training  
27 _TENSOR_AND_DATA_PARALLEL_GROUP = None  
28 # Expert parallel group that the current rank belongs to.  
29 _TENSOR_AND_EXPERT_PARALLEL_GROUP = None  
30 _DATA_MODULO_EXPERT_PARALLEL_GROUP = None  
31 _DATA_MODULO_EXPERT_PARALLEL_GROUP_GLOO = None  
32  
33  
34 _VIRTUAL_PIPELINE_MODEL_PARALLEL_RANK = None  
35 _VIRTUAL_PIPELINE_MODEL_PARALLEL_WORLD_SIZE = None  
36 _PIPELINE_MODEL_PARALLEL_SPLIT_RANK = None  
37  
38 # These values enable us to change the mpu sizes on the fly.  
39 _MPU_TENSOR_MODEL_PARALLEL_WORLD_SIZE = None  
40 _MPU_PIPELINE_MODEL_PARALLEL_WORLD_SIZE = None  
41 _MPU_EXPERT_MODEL_PARALLEL_WORLD_SIZE = None
```

模型进行分组，初始化进程组相关的全局变量

```
97 def initialize_model_parallel(gpus:  
172     GPUs of context parallelism on data parallel group for  
173     weight gradient all-reduce.  
174  
175     nccl_communicator_config_path (str, default = None):  
176         Path to the yaml file of NCCL communicator configurations.  
177         `min_ctas`, `max_ctas`, and `cga_cluster_size` can be set  
178         for each communicator.  
179  
180     Let's say we have a total of 16 GPUs denoted by g0 ... g15 and we  
181     use 2 GPUs to parallelize the model tensor, and 4 GPUs to parallelize  
182     the model pipeline. The present function will  
183     create 8 tensor model-parallel groups, 4 pipeline model-parallel groups  
184     and 8 data-parallel groups as:  
185         8 data_parallel groups:  
186             [g0, g2], [g1, g3], [g4, g6], [g5, g7], [g8, g10], [g9, g11], [g12, g14]  
187         8 tensor model-parallel groups:  
188             [g0, g1], [g2, g3], [g4, g5], [g6, g7], [g8, g9], [g10, g11], [g12, g13]  
189         4 pipeline model-parallel groups:  
190             [g0, g4, g8, g12], [g1, g5, g9, g13], [g2, g6, g10, g14], [g3, g7, g11],  
191             Note that for efficiency, the caller should make sure adjacent ranks  
192             are on the same DGX box. For example if we are using 2 DGX-1 boxes  
193             with a total of 16 GPUs, rank 0 to 7 belong to the first box and  
194             ranks 8 to 15 belong to the second box.  
195  
196     """  
197     # Get world size and rank. Ensure some consistencies.  
198     assert torch.distributed.is_initialized()  
199     world_size: int = torch.distributed.get_world_size()  
200  
201     if (  
202         world_size  
203         % (tensor_model_parallel_size * pipeline_model_parallel_size * context_parallel_size))
```



initialize_model_parallel() 并行配置初始化

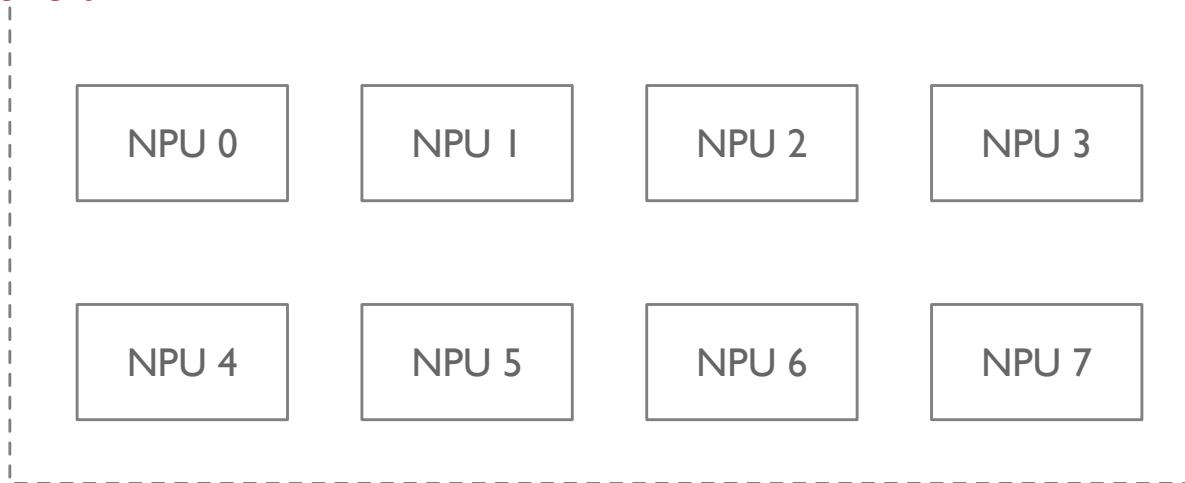
```
negatron > core > parallel_state.py > initialize_model_parallel
97     def initialize_model_parallel():
172         """模型进行分组，初始化进程组相关的全局变量
173             GPUs of context parallelism on data parallel group for
174             weight gradient all-reduce.
175
176             nccl_communicator_config_path (str, default = None):
177                 Path to the yaml file of NCCL communicator configurations.
178                 `min_ctas` , `max_ctas` , and `cga_cluster_size` can be set
179                 for each communicator.
180
181             Let's say we have a total of 16 GPUs denoted by g0 ... g15 and we
182             use 2 GPUs to parallelize the model tensor, and 4 GPUs to parallelize
183             the model pipeline. The present function will
184             create 8 tensor model-parallel groups, 4 pipeline model-parallel groups
185             and 8 data-parallel groups as:
186             8 data_parallel groups:
187                 [g0, g2], [g1, g3], [g4, g6], [g5, g7], [g8, g10], [g9, g11], [g12, g14]
188             8 tensor model-parallel groups:
189                 [g0, g1], [g2, g3], [g4, g5], [g6, g7], [g8, g9], [g10, g11], [g12, g13]
190             4 pipeline model-parallel groups:
191                 [g0, g4, g8, g12], [g1, g5, g9, g13], [g2, g6, g10, g14], [g3, g7, g11]
192             Note that for efficiency, the caller should make sure adjacent ranks
193             are on the same DGX box. For example if we are using 2 DGX-1 boxes
194             with a total of 16 GPUs, rank 0 to 7 belong to the first box and
195             ranks 8 to 15 belong to the second box.
196
197             """
198             # Get world size and rank. Ensure some consistencies.
199             assert torch.distributed.is_initialized()
200             world_size: int = torch.distributed.get_world_size()
201
202             if (
203                 world_size
204                 % (tensor_model_parallel_size * pipeline_model_parallel_size * context_parallel_size)
205             ) != 0:
206                 raise ValueError("world_size must be divisible by tensor_parallel_size, pipeline_parallel_size, and context_parallel_size")
207
208             # Create tensor parallel groups
209             tensor_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % tensor_parallel_size == 0])
210
211             # Create pipeline parallel groups
212             pipeline_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % pipeline_parallel_size == 0])
213
214             # Create data parallel groups
215             data_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % data_parallel_size == 0])
216
217             # Create tensor model-parallel groups
218             tensor_model_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % tensor_parallel_size == 0 and i % pipeline_parallel_size == 0])
219
220             # Create pipeline model-parallel groups
221             pipeline_model_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % tensor_parallel_size == 0 and i % data_parallel_size == 0])
222
223             # Create data-parallel groups
224             data_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % pipeline_parallel_size == 0 and i % data_parallel_size == 0])
225
226             # Create context parallel groups
227             context_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % tensor_parallel_size == 0 and i % pipeline_parallel_size == 0 and i % data_parallel_size == 0])
228
229             # Create weight gradient all-reduce group
230             weight_gradient_all_reduce_group = torch.distributed.new_group(ranks=[i for i in range(world_size)])
231
232             # Create NCCL communicator configuration path
233             nccl_communicator_config_path = None
234
235             # Create NCCL communicator
236             nccl_communicator = NCCLCommunicator(nccl_communicator_config_path)
237
238             # Create tensor parallel group
239             tensor_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % tensor_parallel_size == 0])
240
241             # Create pipeline parallel group
242             pipeline_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % pipeline_parallel_size == 0])
243
244             # Create data parallel group
245             data_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % data_parallel_size == 0])
246
247             # Create tensor model-parallel group
248             tensor_model_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % tensor_parallel_size == 0 and i % pipeline_parallel_size == 0])
249
250             # Create pipeline model-parallel group
251             pipeline_model_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % tensor_parallel_size == 0 and i % data_parallel_size == 0])
252
253             # Create data-parallel group
254             data_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % pipeline_parallel_size == 0 and i % data_parallel_size == 0])
255
256             # Create context parallel group
257             context_parallel_group = torch.distributed.new_group(ranks=[i for i in range(world_size) if i % tensor_parallel_size == 0 and i % pipeline_parallel_size == 0 and i % data_parallel_size == 0])
258
259             # Create weight gradient all-reduce group
260             weight_gradient_all_reduce_group = torch.distributed.new_group(ranks=[i for i in range(world_size)])
261
262             # Create NCCL communicator
263             nccl_communicator = NCCLCommunicator(nccl_communicator_config_path)
```

- 假定16 GPU , 两个 node , rank 0~7 第一个节点 , rank 8 ~15 属于第二个节点 :

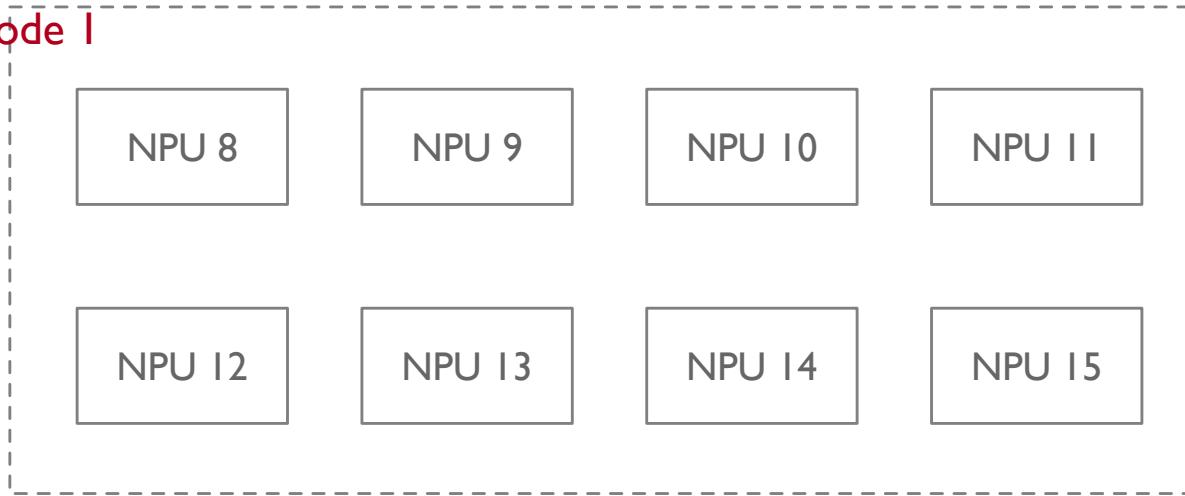
- TP 组大小 2 , 16 个 GPU 被分成 8 组 : [g0, g1], [g2, g3], [g4, g5], [g6, g7], [g8, g9], [g10, g11], [g12, g13], [g14, g15]
- PP 组大小 4 , 16 个 GPU 被分成 4 组 : [g0, g4, g8, g12], [g1, g5, g9, g13], [g2, g6, g10, g14], [g3, g7, g11, g15]
- DP 组大小 2 , 16 个 GPU 被分成 8 组 : [g0, g2], [g1, g3], [g4, g6], [g5, g7], [g8, g10], [g9, g11], [g12, g14], [g13, g15]

并行配置 & 节点情况

Node 0

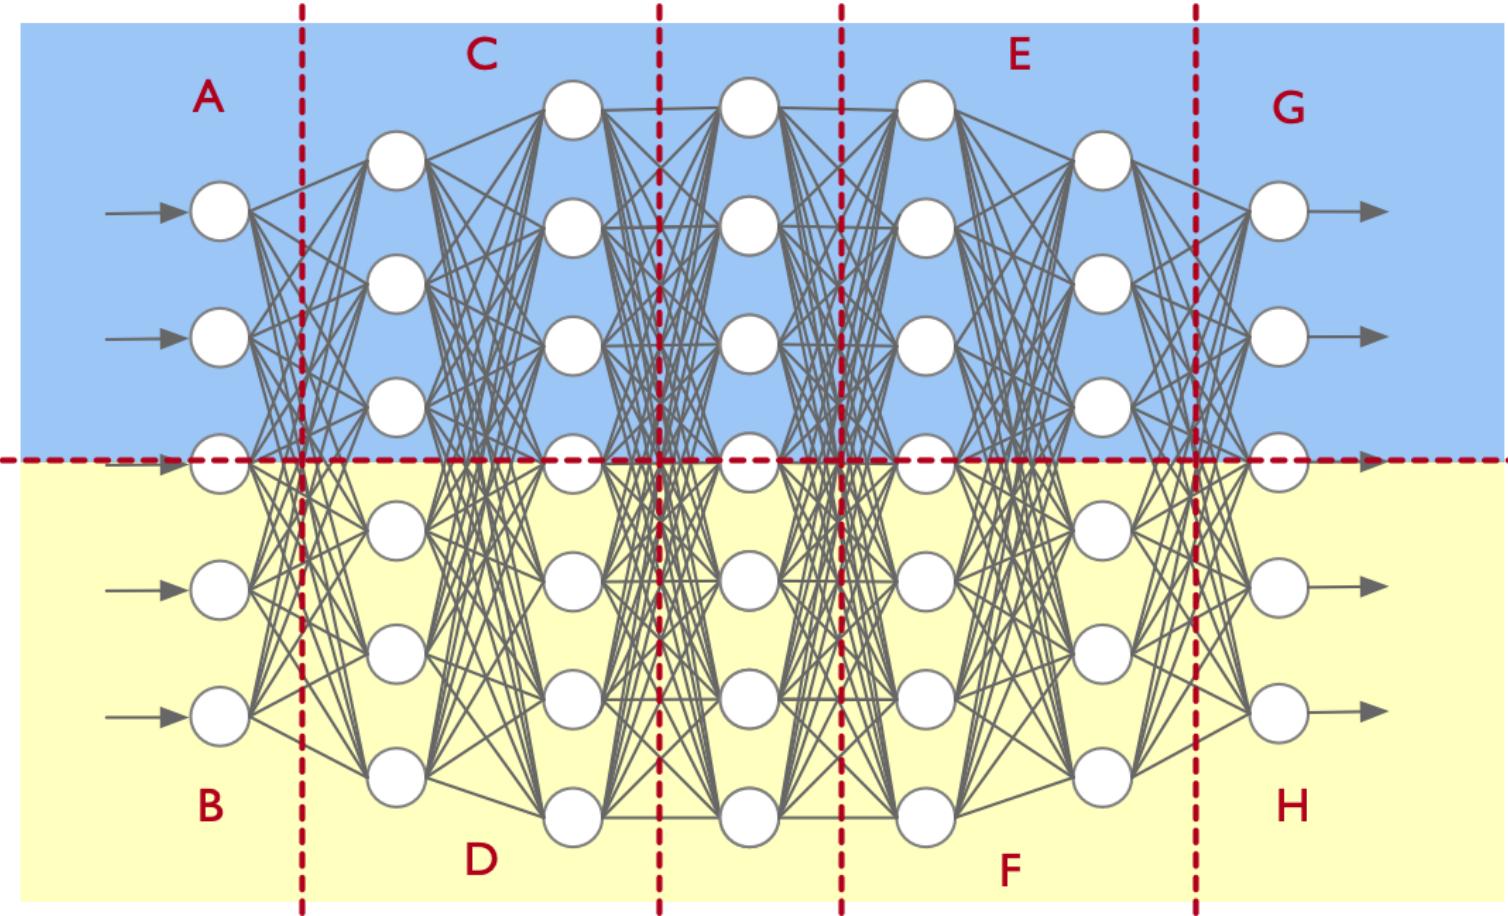


Node 1



- Node 1 : NPU 0 ~ 7
- Node 2 : NPU 8 ~ 15
- world size = 16
- tensor model parallel size = 2
- pipeline model parallel size = 4
- data parallel size = 2

模型并行切分配置



- world size = 16
- tensor model parallel size = 2
- pipeline model parallel size = 4
- 模型并行 (TP+PP) :
 - TP 组大小 2 , 16 个 GPU 被分成 8 组 :
[g0, g1], [g2, g3], [g4, g5], [g6, g7], [g8, g9],
[g10, g11], [g12, g13], [g14, g15]
 - PP 组大小 4 , 16 个 GPU 被分成 4 组 :
[g0, g4, g8, g12], [g1, g5, g9, g13], [g2, g6, g10, g14], [g3, g7, g11, g15]

并行配置 & 处理进程 process groups

```
 2 world_size = 16
 3 tensor_model_parallel_size = 2
 4 pipeline_model_parallel_size = 4
 5 data_parallel_size = world_size // (tensor_model_parallel_size *
 6 | | | | | | | | | | | | | | pipeline_model_parallel_size) * 2
 7 num_tensor_model_parallel_groups = world_size // tensor_model_parallel_size # 8
 8 num_pipeline_model_parallel_groups = world_size // pipeline_model_parallel_size # 4
 9 num_data_parallel_groups = world_size // data_parallel_size # 8
10
11 # Build the data-parallel groups.
12 print("Build DP Groups :")
13 all_data_parallel_group_ranks = []
14 for i in range(pipeline_model_parallel_size):
15     start_rank = i * num_pipeline_model_parallel_groups
16     end_rank = (i + 1) * num_pipeline_model_parallel_groups
17     for j in range(tensor_model_parallel_size):
18         ranks = range(start_rank + j, end_rank,
19 | | | | | | | | | | | | | | tensor_model_parallel_size)
20         all_data_parallel_group_ranks.append(list(ranks))
21 print(all_data_parallel_group_ranks)
```

从 TP 看，分成 8 个进程 group
从 PP 看，分成 4 个进程 group
从 DP 看分成 8 个进程 group
即有 8 个 DDP，每个 DDP 包括 2 个 rank

3.2 Megatron-LM

TP 张量并行

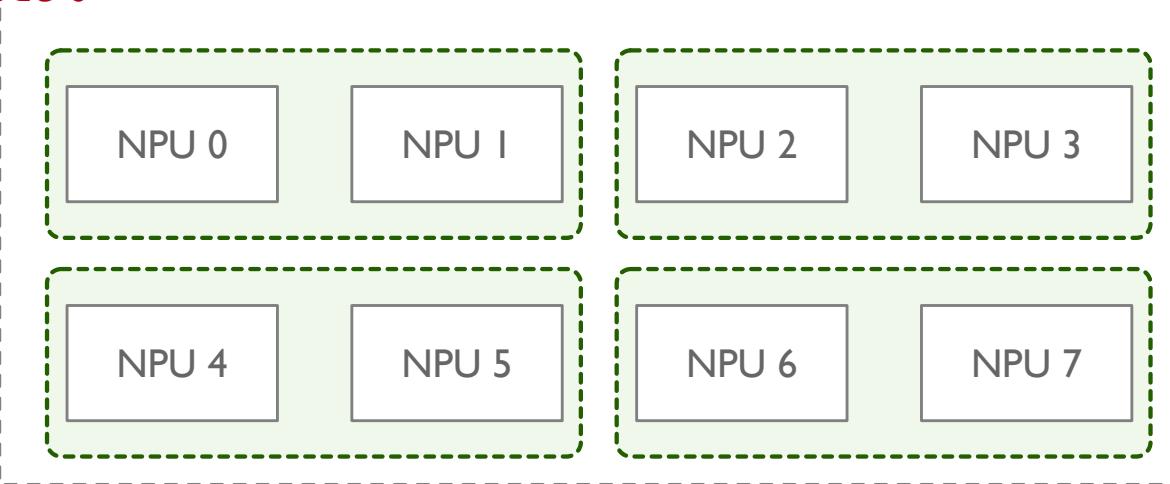
TP 进程组分配

```
350
351     # Build the tensor model-parallel groups.
352     global _TENSOR_MODEL_PARALLEL_GROUP
353     assert (
354         _TENSOR_MODEL_PARALLEL_GROUP is None
355     ), 'tensor model parallel group is already initialized'
356     for i in range(num_tensor_model_parallel_groups): 每一个 tensor-model-parallel group
357         ranks = range(i * tensor_model_parallel_size, 的 rank 一定相邻
358                         (i + 1) * tensor_model_parallel_size)
359         group = torch.distributed.new_group(
360             ranks, pg_options=get_nccl_options('tp', nccl_comm_cfgs)
361         )
362         if rank in ranks:
363             _TENSOR_MODEL_PARALLEL_GROUP = group 记录了当前 rank 的进程组信息
364                                         e.g., rank2 = group([g2, g3])
```

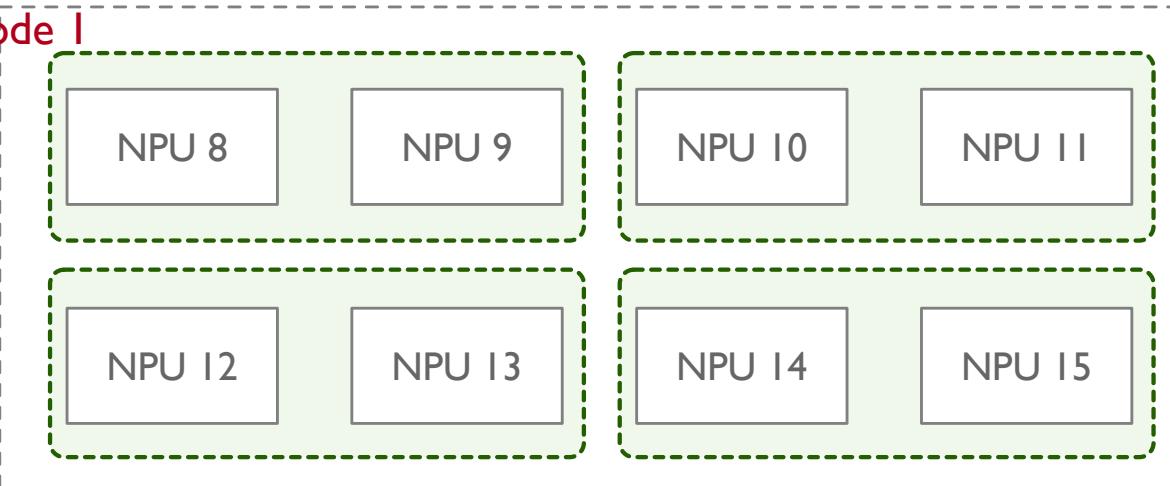


TP 进程组 rank 分配

Node 0



Node 1



- world size = 16 , TP = 2
- group $16 / 2 = 8$
 - 分为 8 个进程组 group , 每个进程组有两个 rank
 - $[g0, g1]$ 为大模型某一层切分为 2 份 , 分别被 rank $g0, g1$ 执行
 - $[g2, g3]$ 为大模型另一层切分为2份 , 分别被 rank $g2, g3$ 执行
 - group 中进行通信主要用到 `reduce()`、`gather()`、`split()`

TP 进程组使用

```
520
521     def get_tensor_model_parallel_group(check_initialized=True):
522         """Get the tensor model parallel group the caller rank belongs to."""
523         if check_initialized:          返回自身 rank 对应的 TP 进程组 group
524             assert (
525                 _TENSOR_MODEL_PARALLEL_GROUP is not None
526             ), 'tensor model parallel group is not initialized'
527         return _TENSOR_MODEL_PARALLEL_GROUP
528
```

```
14
15     def _reduce(input_):
16         """All-reduce the input tensor across model parallel group."""
17
18         # Bypass the function if we are using only 1 GPU.
19         if get_tensor_model_parallel_world_size() == 1:
20             return input_
21
22         # All-reduce.
23         torch.distributed.all_reduce(input_, group=get_tensor_model_parallel_group())
24
25         return input_
26
```

当 TP 反向传播时，利用 _TENSOR_MODEL_PARALLEL_GROUP group 组内进行集合通信

3.2 Megatron-LM

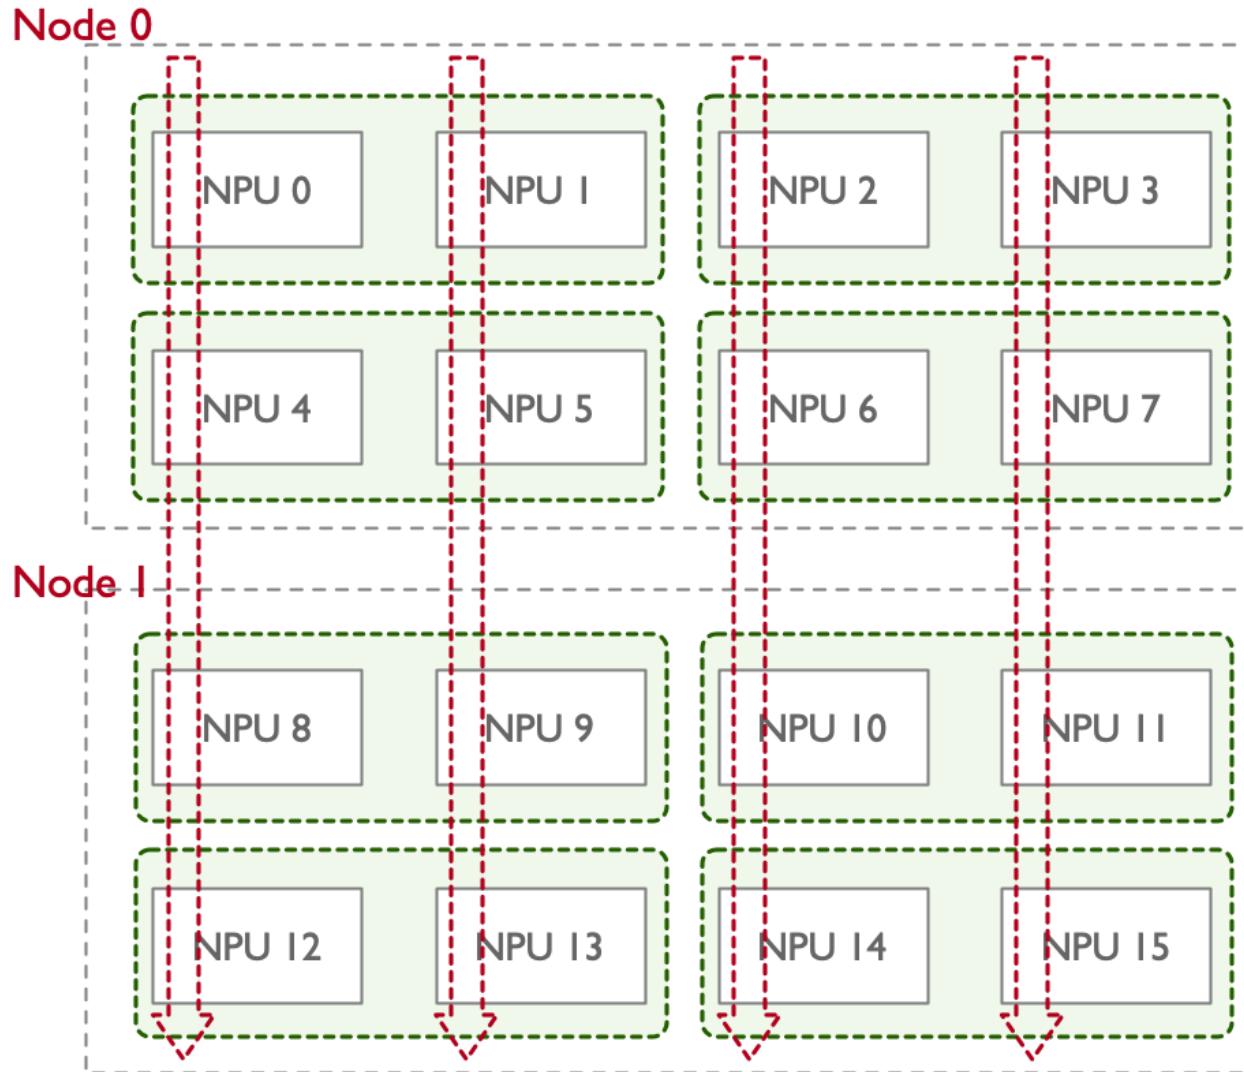
PP 模型并行

PP 进程组分配

```
370     |     _PIPELINE_MODEL_PARALLEL_GROUP is None
371     | ), 'pipeline model parallel group is already initialized'
372     | global _EMBEDDING_GROUP
373     | global _EMBEDDING_GLOBAL_RANKS
374     | assert _EMBEDDING_GROUP is None, 'embedding group is already initialized'
375     | global _POSITION_EMBEDDING_GROUP
376     | global _POSITION_EMBEDDING_GLOBAL_RANKS
377     | assert _POSITION_EMBEDDING_GROUP is None, 'position embedding group is already initialized'
378     | for i in range(num_pipeline_model_parallel_groups):
379     |     ranks = range(i, world_size, num_pipeline_model_parallel_groups) rank 每隔 n // p 个取一个
380     |     group = torch.distributed.new_group(
381     |         ranks, pg_options=get_nccl_options('pp', nccl_comm_cfgs)
382     |     )
383     |     if rank in ranks:
384     |         _PIPELINE_MODEL_PARALLEL_GROUP = group 当前 rank 对应的 PP 进程组
385     |         _PIPELINE_GLOBAL_RANKS = ranks
386     |         # Setup embedding group (to exchange 当前进程组的 ranks
387     |         # first and last stages). e.g., rank2 = [g2, g6, g10, g14]
388     |         if len(ranks) > 1:
```



PP 进程组 rank 分配



- world size = 16 , PP = 4
- group $16 / 4 = 4$
 - 流水线 group 每隔 $n // p$ 个取一个 , stage i 的 rank 范围是 : $[(i-1) * n//p, (i) * n//p]$
 - 分为 4 个进程组 group , 每个进程组有 4 个 rank , 每个 rank 执行 Pipeline 一个 Stage
 - Pipeline Stage 为 4 , 串行执行。即 group [g0, g4, g8, g12] 4 个 NPU 串行。

PP 进程组使用

```
120  def _batched_p2p_ops(
121      *,
122      tensor_send_prev: Optional[torch.Tensor],
123      tensor_recv_prev: Optional[torch.Tensor],
124      tensor_send_next: Optional[torch.Tensor],
125      tensor_recv_next: Optional[torch.Tensor],
126      group: torch.distributed.ProcessGroup
127  ):
128      ops = []
129      if tensor_send_prev is not None:
130          send_prev_op = torch.distributed.P2P0p(
131              torch.distributed.isend,
132              tensor_send_prev,
133              get_pipeline_model_parallel_prev_rank(),
134              group,
135          )
136          ops.append(send_prev_op)
137      if tensor_recv_prev is not None:
138          recv_prev_op = torch.distributed.P2P0p(
139              torch.distributed.irecv,
```

当 TP 进行通信时，使用 P2P0p 点对点的通信方式，而不是集合通信方式



PP 进程组使用

```
529
530     def get_pipeline_model_parallel_group():
531         """Get the pipeline model parallel group the caller rank belongs to."""
532         assert (
533             _PIPELINE_MODEL_PARALLEL_GROUP is not None
534         ), 'pipeline_model parallel group is not initialized'
535         return _PIPELINE_MODEL_PARALLEL_GROUP
536
```

```
862
863     def get_pipeline_model_parallel_next_rank():
864         """Return the global rank that follows the caller in the pipeline"""
865         assert _PIPELINE_GLOBAL_RANKS is not None, "Pipeline parallel group is not initialized"
866         rank_in_pipeline = get_pipeline_model_parallel_rank()
867         world_size = get_pipeline_model_parallel_world_size()
868         return _PIPELINE_GLOBAL_RANKS[(rank_in_pipeline + 1) % world_size]
869
```

通过 `_PIPELINE_GLOBAL_RANKS` 得到了进程组内 ranks 信息

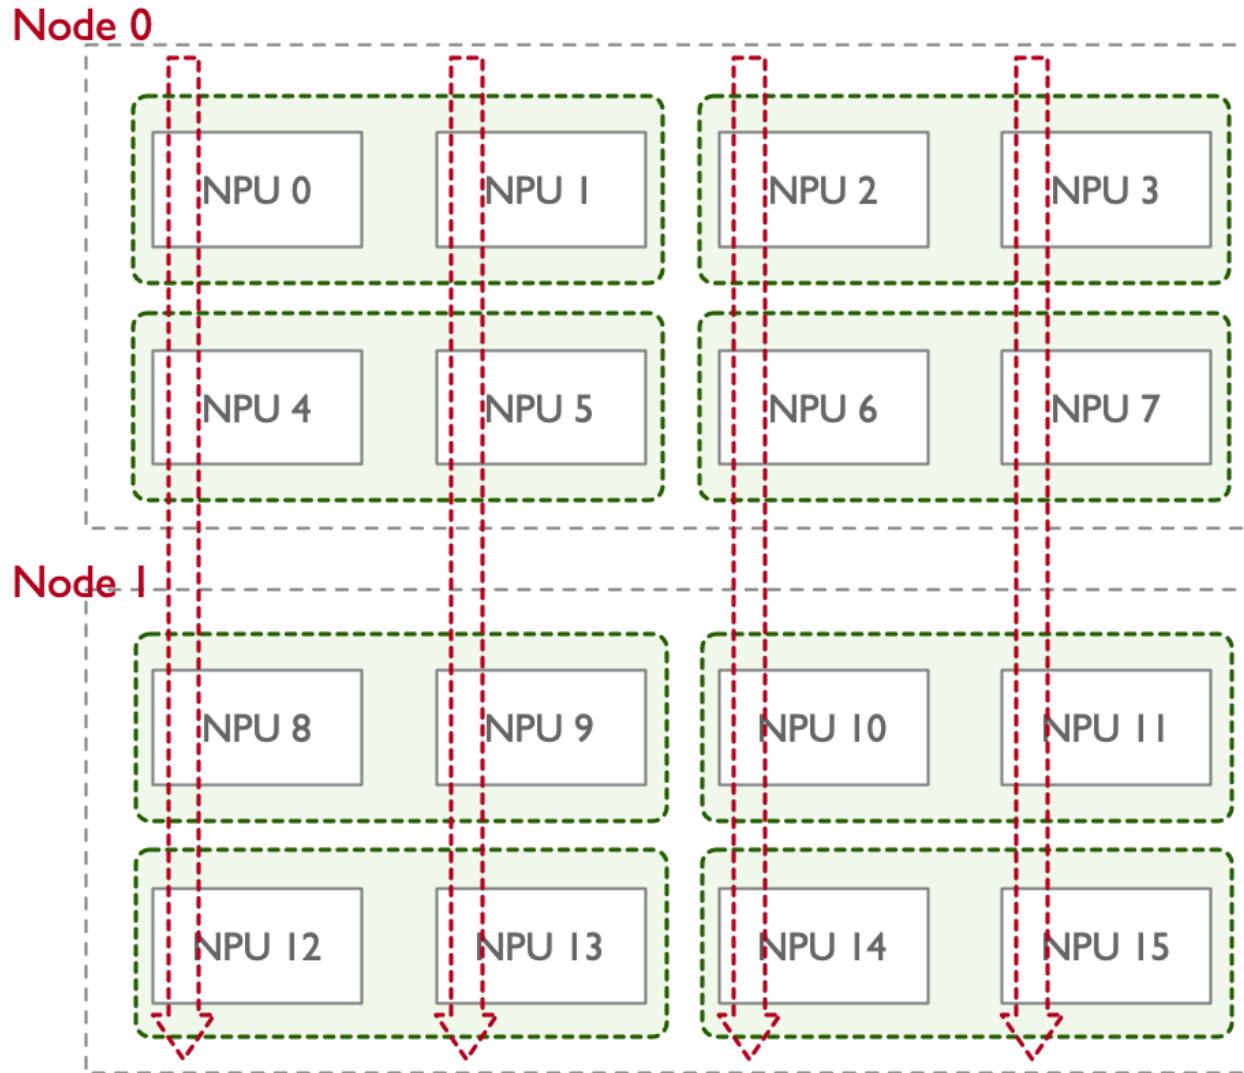
```
855     def get_pipeline_model_parallel_last_rank():
856         """Return the global rank of the last process in the pipeline for the
857         current tensor parallel group"""
858         assert _PIPELINE_GLOBAL_RANKS is not None, "Pipeline parallel group is not initialized"
859         last_rank_local = get_pipeline_model_parallel_world_size() - 1
860         return _PIPELINE_GLOBAL_RANKS[last_rank_local]
```



3.2 Megatron-LM

DP 数据并行

DP 进程组 rank 分配



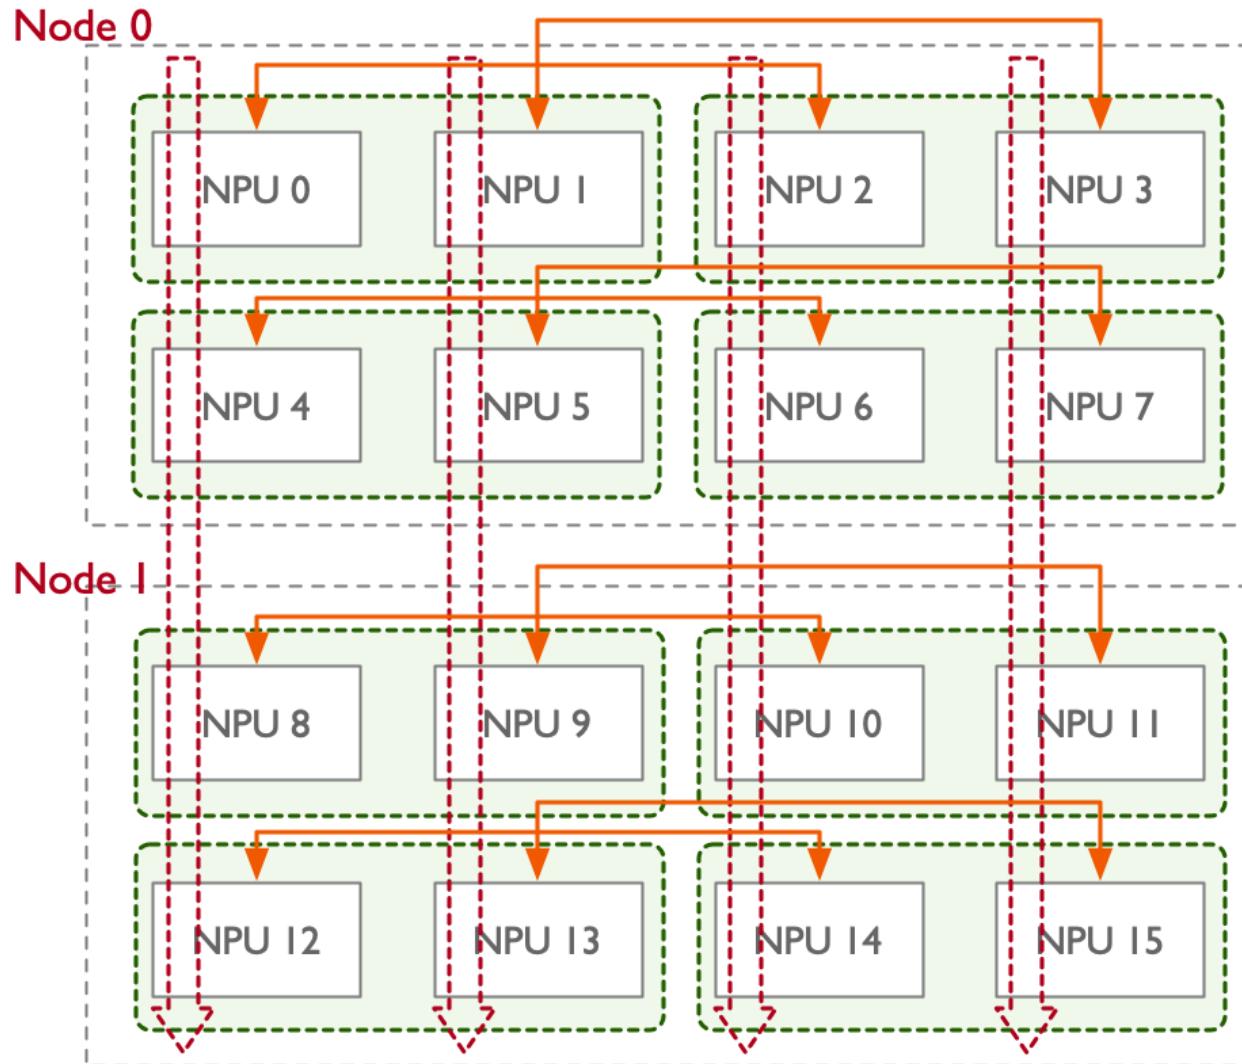
- $T \times P$ 即放下一个大模型所需要的 NPU 资源数：
 - $d = (\text{总 NPU} / \text{一个模型需要资源}) = n / (t * p)$
 - n 个 NPU 可以同时训练 d 个大模型，可用 d 个 mini-batches 输入到 d 个大模型中进行训练
 - 因此 → 数据并行 DP 维度为 d

DP 进程组分配

```
264     global _DATA_PARALLEL_GLOBAL_RANKS_WITH_CP
265     assert _DATA_PARALLEL_GROUP is None, 'data parallel group is already initialized'
266     all_data_parallel_group_ranks_with_cp = []
267     for i in range(pipeline_model_parallel_size): |. 遍历流水线深度
268         start_rank = i * num_pipeline_model_parallel_groups
269         end_rank = (i + 1) * num_pipeline_model_parallel_groups
270         for j in range(context_parallel_size * tensor_model_parallel_size):
271             ranks = range(
272                 start_rank + j, end_rank, context_parallel_size * tensor_model_parallel_size
273             )
274             group = torch.distributed.new_group(
275                 ranks, pg_options=get_nccl_options('dp', nccl_comm_cfgs)
276             )
277             group_gloo = torch.distributed.new_group(ranks, backend="gloo")
278             if rank in ranks:
279                 _DATA_PARALLEL_GROUP = group |. 当前 rank 对应的 DP 进程组
280                 _DATA_PARALLEL_GROUP_GLOO = group_gloo
281                 _DATA_PARALLEL_GLOBAL_RANKS = ranks |. 当前 DP 进程组的 ranks se.g.,rank2 = [g1,g3]
282             for j in range(tensor_model_parallel_size):
283                 ranks_with_cp = range(start_rank + j, end_rank, tensor_model_parallel_size)
284                 all_data_parallel_group_ranks_with_cp.append(list(ranks_with_cp))
285                 group_with_cp = torch.distributed.new_group(
286                     ranks_with_cp, pg_options=get_nccl_options('dp_cp', nccl_comm_cfgs)
287                 )
```



DP 进程组 rank 分配



- world size = 16 , PP = 4 , TP = 2
- DP group $16 / (4 * 2) = 8$
 - 流水线分成 P 个 stage，每个 stage 有 $n // p$ 个 NPU，stage i 的 rank 范围： $[i * n//p, (i+1) * n//p]$ 。e.g. rank 2 rank [0, 1, 2, 3]。
 - 每个 stage 中 ranks，每隔 t 取一个作为数据并行 group 中一份 → 每 DP group 大小为 $n // p // t = d$
 - 通常在 all_reduce() 时候用到

DP 进程组使用

```
538     def get_data_parallel_group(with_context_parallel=False):
539         """Get the data parallel group the caller rank belongs to."""
540         if with_context_parallel:
541             assert (
542                 _DATA_PARALLEL_GROUP_WITH_CP is not None
543             ), 'data parallel group with context parallel combined is not initialized'
544             return _DATA_PARALLEL_GROUP_WITH_CP
```

```
96     def average_losses_across_data_parallel_group(losses):
97         """Reduce a tensor of losses across all GPUs."""
98         averaged_losses = torch.cat(
99             [loss.clone().detach().view(1) for loss in losses])
100        torch.distributed.all_reduce(averaged_losses,
101                                     group=mpu.get_data_parallel_group())
102        averaged_losses = averaged_losses / \
103            torch.distributed.get_world_size(group=mpu.get_data_parallel_group())
104        return averaged_losses
```



DP 进程组使用

```
409     if wrap_with_ddp:
410         config = get_model_config(model[0])
411         model = [DDP(config,
412                      model_chunk,
413                      data_parallel_group=mpu.get_data_parallel_group(with_context_parallel=True),
414                      expert_data_parallel_group=mpu.get_data_modulo_expert_parallel_group(),
415                      accumulate_allreduce_grads_in_fp32=args.accumulate_allreduce_grads_in_fp32,
416                      overlap_grad_reduce=args.overlap_grad_reduce,
417                      use_distributed_optimizer=args.use_distributed_optimizer,
418                      # Turn off bucketing for model_chunk 2 onwards, since communication for these
419                      # model chunks is overlapped with compute anyway.
420                      disable_bucketing=(model_chunk_idx > 0),
421                      check_for_nan_in_grad=args.check_for_nan_in_loss_and_grad)
422         for (model_chunk_idx, model_chunk) in enumerate(model)]
```

```
238     if torch.distributed.is_initialized() and \
239         mpu.get_data_parallel_world_size() > 1 and \
240         args.data_parallel_random_init:
241         rng_state_list = \
242             [None for i in range(mpu.get_data_parallel_world_size())]
243         torch.distributed.all_gather_object(
244             rng_state_list,
245             rng_state,
246             group=mpu.get_data_parallel_group())
247     else:
```



3.2 Megatron-LM

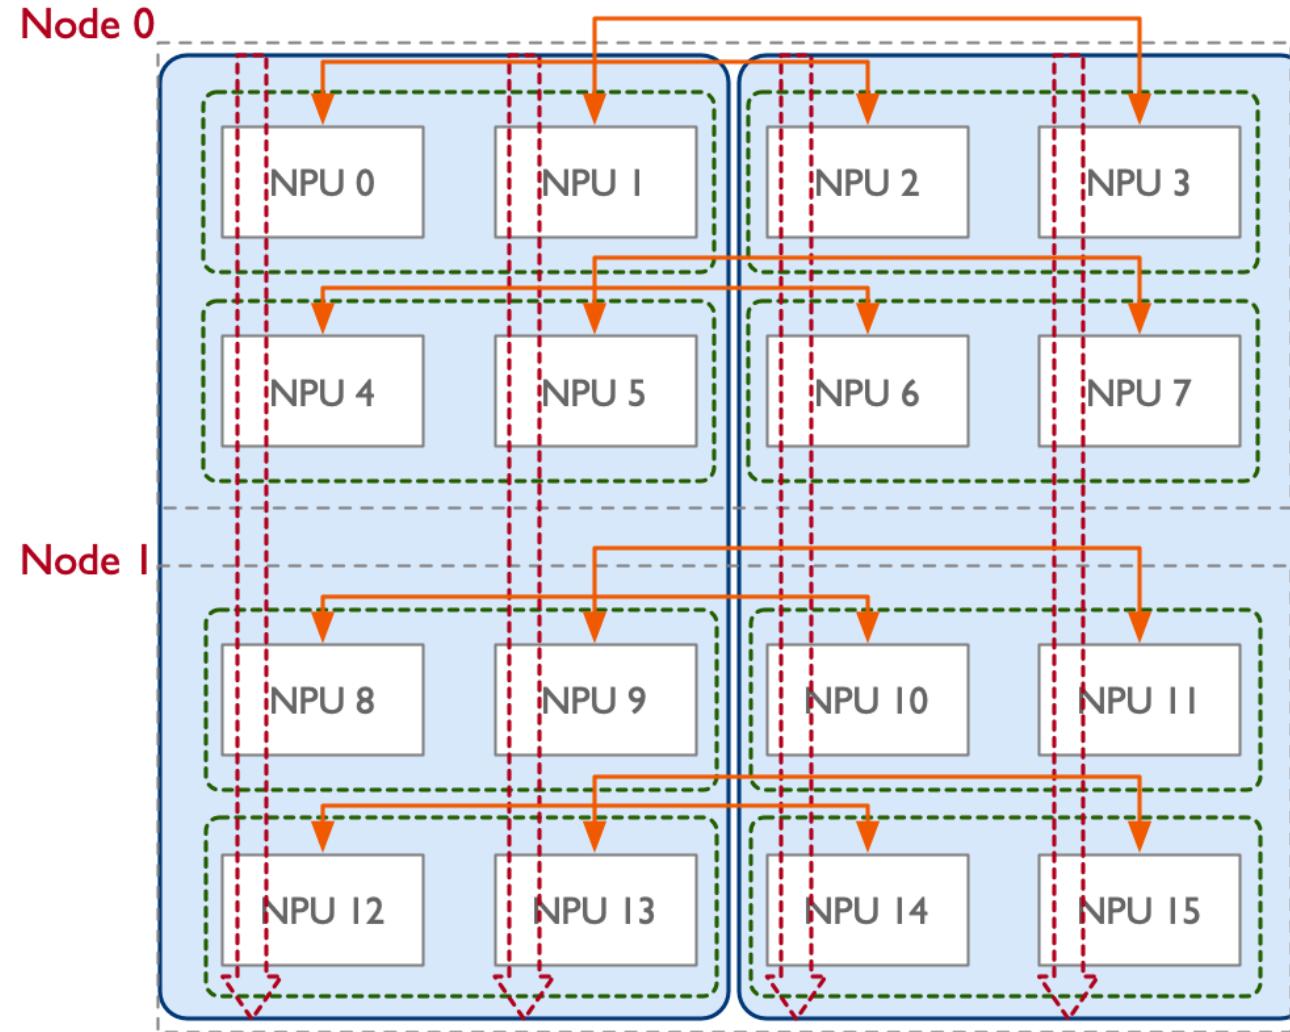
多维并行与配置关系

TPxPP 进程组 rank 分配

```
336
337     # Build the model-parallel groups.
338     global _MODEL_PARALLEL_GROUP
339     assert _MODEL_PARALLEL_GROUP is None, 'model parallel group is already initialized'
340     for i in range(data_parallel_size * context_parallel_size):
341         ranks = [
342             data_parallel_group_ranks_with_cp[i]
343             for data_parallel_group_ranks_with_cp in all_data_parallel_group_ranks_with_cp
344         ]
345         group = torch.distributed.new_group(
346             ranks, pg_options=get_nccl_options('mp', nccl_comm_cfgs)
347         )
348         if rank in ranks:          得到当前 rank 对应的模型 group
349             _MODEL_PARALLEL_GROUP = group e.g., group 1 = [2, 3, 6, 7, 10, 11, 14, 15]
350
```



TPxPP 进程组 rank 分配



- world size = 16 , PP = 4 , TP = 2
- DP group $16 / (4*2) = 8$
 - 从 DP 域中获取第 i 个 rank 作为 MP 的 Group 中的 rank
 - 一共有 DP 个模型进程组，即对应 TP xPP 为一个大模型所需资源

PTD 并行 demo

```
 2 world_size = 16
 3 tensor_model_parallel_size = 2
 4 pipeline_model_parallel_size = 4
 5 data_parallel_size = world_size // (tensor_model_parallel_size *
 6 | | | | | | | | | pipeline_model_parallel_size) # 2
 7 num_tensor_model_parallel_groups = world_size // tensor_model_parallel_size # 8
 8 num_pipeline_model_parallel_groups = world_size // pipeline_model_parallel_size # 4
 9 num_data_parallel_groups = world_size // data_parallel_size # 8
10
11 # Build the data-parallel groups.
12 print("Build DP Groups :")
13 all_data_parallel_group_ranks = []
14 for i in range(pipeline_model_parallel_size):
15     start_rank = i * num_pipeline_model_parallel_groups
16     end_rank = (i + 1) * num_pipeline_model_parallel_groups
17     for j in range(tensor_model_parallel_size):
18         ranks = range(start_rank + j, end_rank,
19 | | | | tensor_model_parallel_size)
20         all_data_parallel_group_ranks.append(list(ranks))
21 print(all_data_parallel_group_ranks)
22
```

PTD 并行 demo

```
22
23     # Build the model-parallel groups.
24     print("__iter__ Group:")
25     for i in range(data_parallel_size):
26         ranks = [data_parallel_group_ranks[i]
27                  | | | for data_parallel_group_ranks in all_data_parallel_group_ranks]
28         print(list(ranks))
29
30     # Build the tensor model-parallel groups.
31     print("Build TP Groups:")
32     for i in range(num_tensor_model_parallel_groups):
33         ranks = range(i * tensor_model_parallel_size,
34                        | | | (i + 1) * tensor_model_parallel_size)
35         print(list(ranks))
36
37     # Build the pipeline model-parallel groups and embedding groups
38     print("Build PP Groups :")
39     for i in range(num_pipeline_model_parallel_groups):
40         ranks = range(i, world_size,
41                        | | | num_pipeline_model_parallel_groups)
42         print(list(ranks))
```



PTD 并行 demo 结果

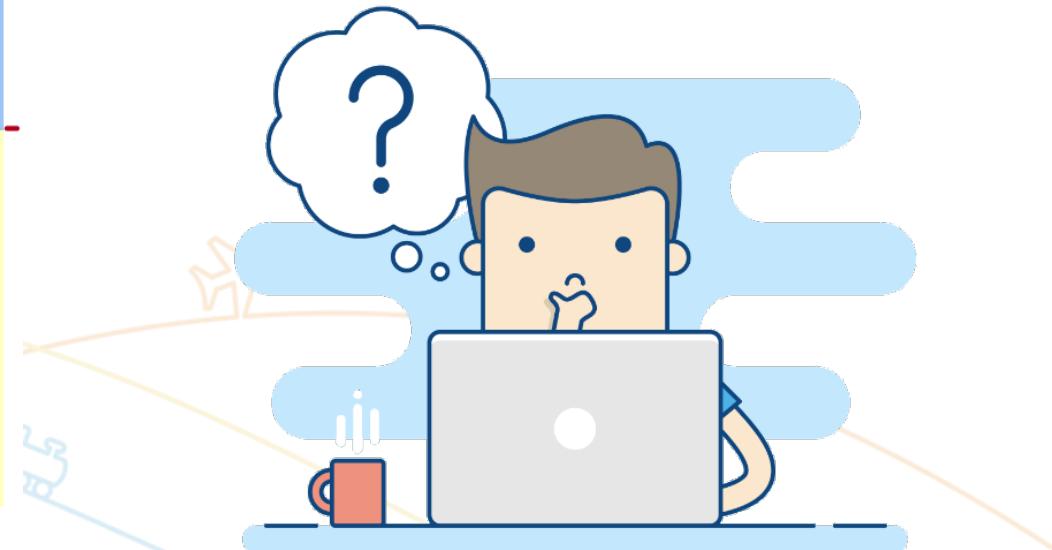
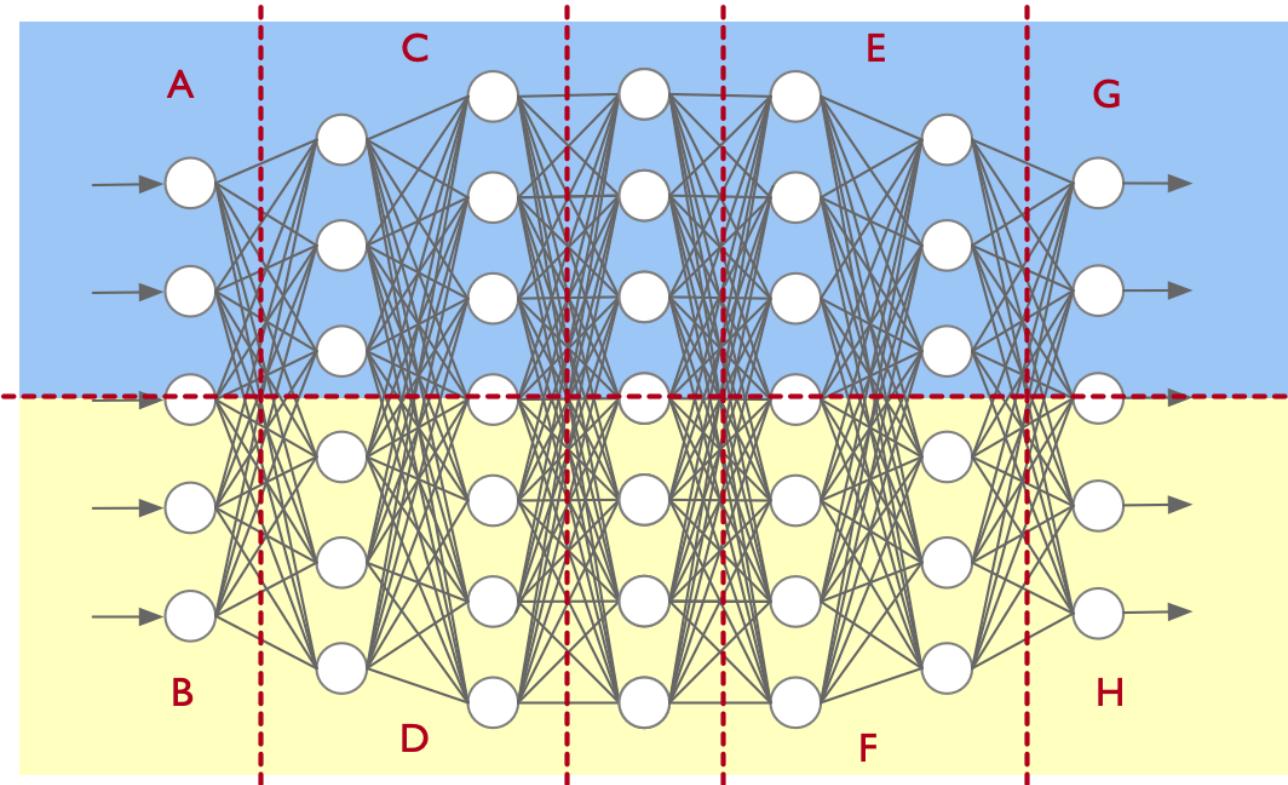
```
1. chenzhongming@chenzomi-2: ~/Workplaces/DLSys_github/06Foundation/07Train/code (zsh)
~/Workplaces/DLSys_github/06Foundation/07Train/code ➜ main • ? ↑1 ⏪1 python ptd_config.py ➜ 21:29:29
Build DP Groups :
[[0, 2], [1, 3], [4, 6], [5, 7], [8, 10], [9, 11], [12, 14], [13, 15]]
Build MP Group:
[0, 1, 4, 5, 8, 9, 12, 13]
[2, 3, 6, 7, 10, 11, 14, 15] 文本
Build TP Groups:
[0, 1]
[2, 3]
[4, 5]
[6, 7]
[8, 9]
[10, 11]
[12, 13]
[14, 15]
Build PP Groups :
[0, 4, 8, 12]
[1, 5, 9, 13]
[2, 6, 10, 14]
[3, 7, 11, 15]
文件夹
code
cover
images
old
演示文稿
01Introduction.pptx
02DeepSpeed.pptx
03Megatron.pptx
04MegatronPTD.pptx
开发者
demo1.py
demo1.sh
demo2.py
ptd_config.py
其他
ds_config.json
世界
world_size = 16
tensor_model_parallel_size = 2
pipeline_model_parallel_size = 4
data_parallel_size = world_size // (tensor_model_parallel_size * pipeline_model_parallel_size) # 2
num_tensor_model_parallel_groups = world_size // tensor_model_parallel_size # 8
num_pipeline_model_parallel_groups =
PPT_Template.pptx
LICENSE
21:29:31
ptd_config.py
Python 脚本 - 2 KB
信息
```

3.2 Megatron-LM

大模型与卡间关系

思考

I. 如何把大模型按照模型层数或者切分好的模块，分块放到对应的 NPU 上？



执行流程

1. rank 根据 PTD 全局变量映射到 NPU
2. 模型初始化通过 offset 根据 rank 生成对应层
3. 模型参数根据 rank 拷贝到对应 NPU 上然后执行训练

1. rank 根据 PTD 全局变量映射到 NPU

```
12 # Intra-layer model parallel group that the current rank belongs to.  
13 _TENSOR_MODEL_PARALLEL_GROUP = None  
14 # Inter-layer model parallel group.  
15 _PIPELINE_MODEL_PARALLEL_GROUP = None  
16 # Model parallel group (both intra and inter layer).  
17 _MODEL_PARALLEL_GROUP = None  
18 # Embedding group.  
19 _EMBEDDING_GROUP = None  
20 # Position embedding group.  
21 _POSITION_EMBEDDING_GROUP = None  
22 # Data parallel group that the current rank belongs to.  
23 _DATA_PARALLEL_GROUP = None  
24 _DATA_PARALLEL_GROUP_GLOO = None  
25 # tensor model parallel group and data parallel group combined  
26 # used for fp8 and moe training  
27 _TENSOR_AND_DATA_PARALLEL_GROUP = None  
28 # Expert parallel group that the current rank belongs to.  
29 _TENSOR_AND_EXPERT_PARALLEL_GROUP = None  
30 _DATA_MODULO_EXPERT_PARALLEL_GROUP = None  
31 _DATA_MODULO_EXPERT_PARALLEL_GROUP_GLOO = None
```

- mpu.initialize_model_parallel() 设置 PTD 等进程组，每个 rank 对应进程都有其全局变量，进程自动映射到 NPU 上：
 - e.g., rank 2 对应进程启动后，从初始化全局变量中获取 data_parallel group: [g0, g2] , tensor model-parallel group: [g2, g3] , pipeline model-parallel group: [g2, g6, g10, g14]。



2. 模型初始化通过 offset 根据 rank 生成对应层

```
1378     class ParallelTransformer(MegatronModule):
1381         def __init__(self, config,
1382                      # With 8 layers, 2 stages, and 2 virtual stages, we want an assignment of
1383                      # layers to stages like (each list is a model chunk):
1384                      # Stage 0: [0, 1]  [4, 5]
1385                      # Stage 1: [2, 3]  [6, 7]
1386                      offset = mpu.get_virtual_pipeline_model_parallel_rank() * (
1387                          config.num_layers // config.virtual_pipeline_model_parallel_size) + \
1388                          (mpu.get_pipeline_model_parallel_rank() * self.num_layers)
1389
1390         else:
1391             # Each stage gets a contiguous set of layers.
1392             if args.model_type == ModelType.encoder_and_decoder and \
1393                 mpu.get_pipeline_model_parallel_world_size() > 1
1394                 pipeline_rank = mpu.get_pipeline_model_parallel_rank
1395                 if layer_type == LayerType.encoder:
1396                     offset = pipeline_rank * self.num_layers
1397                 else:
1398                     num_ranks_in_enc = args.pipeline_model_parallel_split_rank
1399                     offset = (pipeline_rank - num_ranks_in_enc) * self.num_layers
1400
1401             else:
1402                 offset = mpu.get_pipeline_model_parallel_rank() * self.num_layers
```

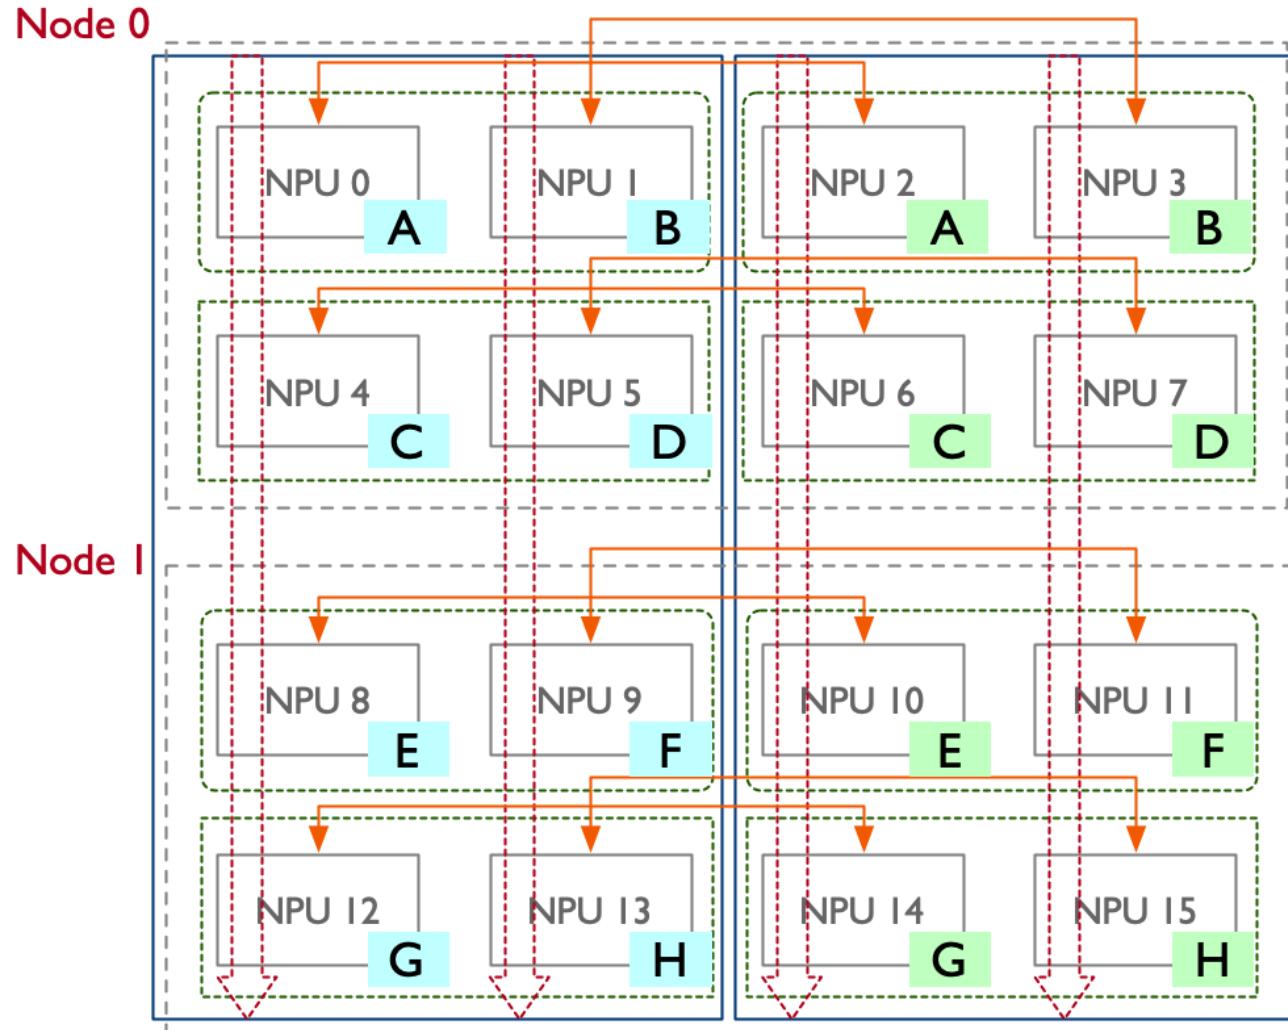
offset 根据 rank 知道自己应该生成模型所属层



2.模型初始化通过 offset 根据 rank 生成对应层

```
1571         # optimizations (e.g., pipeline output deallocation). To remedy
1572         # this, we assign a 'no-op' layer on these ranks, which will
1573         # disconnect the input tensor from the output tensor. 通过 self.layers 来生成对应的层
1574     self.num_layers = 1
1575     self.layers = torch.nn.ModuleList([ NoopTransformerLayer(1) ])
1576 else:
1577     self.layers = torch.nn.ModuleList(
1578         [build_layer(i + 1 + offset) for i in range(self.num_layers)])
```

建立模型与卡间关键



- I. 同名子模块具有同样的网络模型结构与参数，可进行数据并行：
 - e.g., 两个 A 可以数据并行
- I. 纵向层间可进行流水线串行：
 - e.g., A → C → E → G
- I. 横向是流水线的一个stage：
 - e.g., 从 0 开始，相邻为 A & B 为 TP

总结与思考

PTD 并行配置

1. TP 张量并行被用于 intra-node transformer 层：
 - TP 张量并行计算密集且耗费大量带宽，节点内利用高带宽可以高效运行。
2. PP 流水并行主要被用于 inter-node transformer 层：
 - PP 通信带宽占用少，其可以有效利用集群中多卡设计
3. DP 在 PP 和 TP 基础之上进行加持，使得训练可以扩展到更大规模和更快的速度：
 - 尽管 DP 可高效扩展，但不能单独使用 DP 来训练超大模型，a) HBM 不足，b) 数据并行扩展限制

PTD 并行配置

Number of parameters (billion)	Attention heads	Hidden size	Number of layers	Tensor model-parallel size	Pipeline model-parallel size	Number of GPUs	Batch size	Achieved teraFLOP/s per GPU	Percentage of theoretical peak FLOP/s	Achieved aggregate petaFLOP/s
1.7	24	2304	24	1	1	32	512	137	44%	4.4
3.6	32	3072	30	2	1	64	512	138	44%	8.8
7.5	32	4096	36	4	1	128	512	142	46%	18.2
18.4	48	6144	40	8	1	256	1024	135	43%	34.6
39.1	64	8192	48	8	2	512	1536	138	44%	70.8
76.1	80	10240	60	8	4	1024	1792	140	45%	143.8
145.6	96	12288	80	8	8	1536	2304	148	47%	227.1
310.1	128	16384	96	8	16	1920	2160	155	50%	297.4
529.6	128	20480	105	8	35	2520	2520	163	52%	410.2
1008.0	160	25600	128	8	64	3072	3072	163	52%	502.0

Table 1: Weak-scaling throughput for GPT models ranging from 1 billion to 1 trillion parameters.



把AI系统带入每个开发者、每个家庭、
每个组织，构建万物互联的智能世界

Bring AI System to every person, home and
organization for a fully connected,
intelligent world.

Copyright © 2023 XXX Technologies Co., Ltd.
All Rights Reserved.

The information in this document may contain predictive statements including, without limitation, statements regarding the future financial and operating results, future product portfolio, new technology, etc. There are a number of factors that could cause actual results and developments to differ materially from those expressed or implied in the predictive statements. Therefore, such information is provided for reference purpose only and constitutes neither an offer nor an acceptance. XXX may change the information at any time without notice.



Course chenzomi12.github.io

GitHub github.com/chenzomi12/DeepLearningSystem